

Total Cost of Ownership: Hedge Fund AI Consolidation

Industry: Financial Services — Hedge Funds, Asset Managers **Topic:** Platform Consolidation & Cost Reduction **Platform:** Roundtable by Foxtrot Communications

Executive Summary

A typical 100-person hedge fund spends **\$1.3M+ annually** on a patchwork of disconnected AI tools — per-seat ChatGPT licenses, standalone compliance platforms, duplicate cloud API subscriptions, and 2-3 engineers maintaining custom integrations between them. None of these systems talk to each other, and there is no unified audit trail.

Roundtable replaces this entire stack with a single, self-hosted platform at an infrastructure cost of **~\$30K/year** plus licensing. The savings come from eliminating per-seat AI licensing, consolidating duplicate subscriptions, removing integration engineering, and replacing standalone compliance AI platforms — while adding capabilities none of the individual tools provide.

The Problem: AI Sprawl

Every desk picks their own AI. Every team has their own API keys. Every tool has its own data pipeline to the same underlying data.

Trading Desk:	ChatGPT Enterprise + Bloomberg AI	\$180K/yr
Quant Research:	Claude Teams + Custom ML Platform	\$324K/yr
Compliance:	Behavox + NICE Actimize	\$200K/yr
Risk Management:	Internal Python + OpenAI API	\$85K/yr
Operations:	ChatGPT Enterprise	\$36K/yr
Integration Glue:	2-3 Engineers keeping it connected	\$400K/yr
Duplicate Infra:	Each tool's own ETL + storage	\$100K/yr
Total:		~\$1.3M+

And none of it talks to each other. When compliance needs trading data for an investigation, they send an email. When quant builds a risk model, they export a CSV. When a regulator asks for AI interaction records, three different teams scramble to produce logs in three different formats.

The Solution: One Platform, N Workspaces

```
GKE Cluster
├── rt-compliance      Bridge    rt-trading
├── rt-risk            Bridge    rt-compliance
├── rt-quant
├── rt-ops
├──
├── Shared PostgreSQL (all messages, all audit trails)
├── Shared Vertex AI endpoint (usage-based pricing)
├── Shared BigQuery connection (configured once)
├── Shared A2A Agent Pool
│   ├── Bloomberg Agent
│   ├── Kensho Agent
│   └── Custom Risk Model Agent
```

Each workspace is an isolated container with its own AI configuration, tools, and system prompt — but shares infrastructure with every other workspace. Users connect to their team's workspace via browser. No desktop software. No per-seat licensing.

Cost Comparison

Current State: Disconnected Tools

Category	Components	Annual Cost
Per-seat AI licensing	ChatGPT Enterprise (100 seats), Claude Teams (20 seats)	\$96,000
Compliance AI platform	Behavox, NICE Actimize, or equivalent	\$200,000

Custom ML/AI platform	Quant team infrastructure, GPU instances	\$300,000
Cloud AI API accounts	OpenAI, Anthropic, Google — per-team accounts	\$60,000
Integration engineering	2-3 engineers maintaining connectors	\$400,000
Duplicate data pipelines	Each tool's own ETL to the same data	\$100,000
Bloomberg AI add-ons	Per-terminal AI features	\$150,000
	Total	\$1,306,000

Roundtable: Consolidated Platform

Category	Components	Annual Cost
GKE cluster	5 workspace pods, e2-standard-2, spot instances	\$3,000
Cloud SQL	PostgreSQL (db-f1-micro), shared across all workspaces	\$1,200
Vertex AI inference	Gemini Flash + Pro, usage-based (no per-seat)	\$18,000
BigQuery	On-demand pricing, already provisioned	\$2,000
A2A agent hosting	Bloomberg/Kensho adapters, lightweight containers	\$5,000
Roundtable Enterprise license	Cross-workspace bridges, audit, fleet management	TBD

Integration engineering	Agents plug in via A2A — no custom glue	\$0
	Total	~\$30,000 + license

Savings Breakdown

Savings Category	Amount	How
Eliminate per-seat licensing	\$96,000	Workspace-based, not user-based
Eliminate compliance AI platform	\$200,000	Native audit trail in PostgreSQL
Eliminate integration engineering	\$400,000	Tools and agents are shared infrastructure
Consolidate AI API accounts	\$60,000	One Vertex AI endpoint, volume pricing
Reduce custom ML platform	\$200,000	Quant models exposed as A2A agents
Eliminate duplicate pipelines	\$100,000	BigQuery configured once for all workspaces
	Total savings	~\$1,056,000/yr

Why Per-Seat Licensing Is Obsolete

ChatGPT Enterprise charges \$60/user/month regardless of usage. A trader who asks one question a week pays the same as an analyst running 50 queries a day.

Roundtable's model:

- **Infrastructure cost:** Fixed per workspace (~\$600/yr per pod)

- **AI inference cost:** Per token consumed — pay for what you use
- **Users:** Unlimited per workspace at no additional cost

Model	Cost for 100 users
ChatGPT Enterprise	\$72,000/yr
Roundtable + Gemini Flash	~\$6,000/yr

The difference widens with headcount. At 200 users, ChatGPT doubles to \$144K. Roundtable stays at ~\$6K because adding users to an existing workspace costs nothing.

Why Integration Engineering Disappears

Today a typical fund employs 2-3 engineers (\$150-200K each) whose primary job is connecting AI tools to each other and to internal data:

- Building a pipeline from Bloomberg Terminal to the compliance AI platform
- Writing exporters from ChatGPT conversation logs to the regulatory archive
- Maintaining a bridge between the quant team's Python models and the risk dashboard
- Creating Slack bots that forward AI outputs between teams

In Roundtable:

- **Bloomberg** is an A2A agent deployed once, accessible to any workspace
- **BigQuery** is a built-in tool configured once in the ConfigMap
- **Conversation logs** are already in PostgreSQL — no export needed
- **Cross-team AI queries** are native via workspace bridges

The 2-3 integration engineers can be redeployed to higher-value work.

Why Standalone Compliance Platforms Are Redundant

Behavox and NICE Actimize charge \$200K+ for:

Capability	Behavox (\$200K+)	Roundtable (included)
Record all AI interactions	Yes	Yes — <code>messages</code> table
Attribute to user + timestamp	Yes	Yes — <code>user_id</code> , <code>created_at</code>
Search interaction history	Yes	Yes — SQL queries
Cross-desk audit trail	Yes	Yes — <code>source_workspace_id</code>
Tool call logging	No	Yes — every tool call recorded
AI-powered investigation	No	Yes — dedicated compliance AI

Model Arbitrage: Right-Size AI Spend

Each workspace can use a different model — pay for capability only when needed:

Workspace	Model	Cost/1M tokens
Operations	Gemini 2.0 Flash	\$0.075
Compliance (routine)	Gemini 2.0 Flash	\$0.075
Compliance (escalated)	Gemini 2.5 Pro	\$1.25
Quant Research	Gemini 2.5 Pro	\$1.25
Trading Desk	Gemini 2.0 Flash	\$0.075

80% of queries use Flash. Only complex reasoning needs Pro. Per-seat licensing doesn't give you this flexibility.

Beyond Cost: New Capabilities

Cross-Workspace AI Routing

A compliance analyst types `@trading-ai Show me AAPL trades from May 14` — and the trading workspace's AI responds with Bloomberg data. Both workspaces see the exchange. Full audit trail. No email, no Slack, no CSV.

No combination of ChatGPT + Behavox + Bloomberg can do this.

Unified Institutional Memory

Every AI conversation across every desk is in one database. When an analyst leaves, their investigation history stays. When a regulator asks "what AI-assisted decisions led to this trade?" — the answer is a SQL query, not a six-week forensic exercise.

One-Command Desk Deployment

```
GCP_PROJECT=fund-infra ./deploy-gke.sh newdesk "New Desk"
```

A new team gets full AI capabilities in 5 minutes. Today it takes weeks.

Implementation Timeline

Phase	Duration	Deliverable
Pilot	2 weeks	1 workspace, 5-10 users
Validation	2 weeks	Confirm audit trail, test bridges
Rollout	2 weeks	All desks, A2A agents
Decommission	4 weeks	Phase out legacy tools
Total	10 weeks	Full migration

Summary

Metric	Current	Roundtable
--------	---------	------------

Annual AI spend	\$1.3M+	~\$30K + license
Systems to maintain	7+	1
Unified audit trail	No	Yes
Cross-desk AI collaboration	No	Yes
Time to onboard a desk	Weeks	Minutes
Per-seat cost	\$60/user/mo	\$0
Vendor lock-in	High	None (Apache 2.0)
Integration engineers	2-3	0

Roundtable by [Foxtrot Communications](#) — Replace your AI sprawl with one platform your compliance team will actually love.